



E0-261

Database Management Systems

Project #08



Akash Maji [24212]

Utkarsh Sharma [24116]



Agenda


1. About DuckDB
2. Project Overview
3. DuckDB Architecture
4. Execution Model
5. New Join Operator
6. New GroupJoin Operator
7. Test & Experiments
8. Summary & Future Work

Let's Dive In!



About DuckDB

- DuckDB is an open-source analytical (OLAP) database system.
- It runs as a **single-process** without needing server connection.
- The entire database can sit in a *file* with **.db** or **.duckdb** extension.



Features:

- Fast, Reliable, Portable, Open-Source
- Analytical Support
- Columnar, Compressed, Block-Based Storage

In this project,

- we aim to explore the [interface](#) provided by DuckDB,
- and implement certain [operators](#).

1

Firstly, we will discuss the implementation of a new ‘Join’ operator.

2

Secondly, we change the [query planner](#) and [optimizer](#) to pick our join operator.

- whenever a [predetermined](#) condition is met.

3

Thirdly, we implement a new ‘GroupJoin’ operator.

- which will be invoked if a query [containing a ‘Join’ followed by a ‘Group By’](#) is handed to the engine.

DuckDB Architecture

- ***SQLStatement***: query statement in DuckDB is represented as this
- ***Binder***: responsible for **binding** tables and columns to actual physical tables and columns in the catalog.
- ***Planner***: has a ***CreatePlan()*** method that creates the **logical plan tree** from the **AST**.
- The choice of actual algorithm is left to the optimizer based on **estimated cardinality**.
- A **Value** is a unit that holds a **single** value of arbitrary type.
- A **Vector** is the **smallest unit** of data handling holding values in a column.
- A **DataChunk** is a set of Vectors serving as the unit of data processing in the pipeline through operators.
- A **Vector** can hold up to a fixed number of values, defined by **STANDARD_VECTOR_SIZE** (power of 2)

Table	
101	25
102	23
103	31
104	28
105	30
106	27

Chunk1		Chunk2	
101	25	105	30
102	23	106	27
103	31		
104	28		






Execution Model

Push-Based, Pipelined

- Operators **process** the data (e.g. Join) and **push** the data to the next operator (e.g. filter) in the pipeline.
- A long **execution sequence** is broken into multiple **pipeline events** and executed possibly in *parallel*.

Each pipeline has *three* interfaces: **Source**, **Operator** and **Sink**.

```
SELECT stud.sid, enrol.cid FROM stud JOIN enrol  
ON stud.sid = enrol.sid WHERE stud.sage > 25;
```






New Join Operator (AM_US_JOIN)

- **Operator Definition:** The "AM_US_JOIN" physical operator is associated with the "LOGICAL_COMPARISON_JOIN" logical operator and functions as a nested loop join.
- **Operator Selection:** DuckDB's query planner selects join operators based on the estimated cardinality of tuples in participating relations.
- **Threshold Configuration:** The client configuration is tweaked to set a cardinality threshold (e.g., 100 tuples), which is tunable as needed.
- **Physical Plan Generator:** A cardinality check is added to the physical plan generator to ensure input relations are within the threshold.
- **Automatic Selection:** The optimizer automatically chooses "AM_US_JOIN" for small relations when detected as child nodes in the physical plan.

Working (Contd)

```
if(can_do_physical_amus_join(client_config, left, right)){
    plan = make_uniq<PhysicalAmUsJoin>(op, std::move(left), std::move(right), std::move(op.conditions),
                                         op.join_type, op.estimated_cardinality);
    std::cout << "AM_US_JOIN Physical Plan Taken\n";
    return plan;
}
```

```
bool can_do_physical_amus_join(ClientConfig& client_config,
                                LogicalComparisonJoin& op,
                                unique_ptr<PhysicalOperator>& left,
                                unique_ptr<PhysicalOperator>& right){
    bool can_use = false;
    if(PhysicalAmUsJoin::IsSupported(op.conditions, op.join_type)){
        if (left->estimated_cardinality <= client_config.am_us_join_threshold &
            right->estimated_cardinality <= client_config.am_us_join_threshold) {
            can_use = true;
        }
    }
    return can_use;
}
```



Working

- ***PhysicalAmUsJoin*** class extending the ***PhysicalComparisonJoin***
- Actual join is performed by the ***ResolveComplexJoin()*** method by calling ***Execute()*** internally.
- The join starts with getting the left *datachunk*, and operating it will all *datachunks* on the right.
- Then we go to next *datachunk* in left and continue the process until all *datachunks* are exhausted in left.
- For each pair of left and right *datachunk*, there are tuple markers ‘left’ and ‘right’ which progressively move during operation, and the rows are marked first which matches the Join condition.
- The matching positions are kept in a *match_vector*.
- The actual join happens between the left and right *datachunks* using the ***Perform()*** method.
- If matches are **found**, the matched tuples are **sliced** from the input and stored in the output *datachunk*.

Working

- *PhysicalAmUsJoin* class extending the *PhysicalComparisonJoin*
- Actual join is performed by the *ResolveComplexJoin()* method by calling *Execute()* internally.
- The join starts with  "Talk is cheap. Show me the code."
- Then we go to next *datachunk* on the right.
- For each pair of left *datachunk* and right *datachunk* which progressively move during the join.
- The matching position is found in left *datachunk*.
- The actual join happens between the left and right *datachunks* using the *Perform()* method.
- If matches are found, the matched tuples are sliced from the input and stored in the output *datachunk*.

ResolveComplexJoin Algorithm

Algorithm



```

Algorithm ResolveComplexJoin(context, input, chunk, state){

    // Cast state and global state to specific types
    state = cast state to PhysicalAmUsJoinState
    gstate = cast sink_state to AmUsJoinGlobalState

    // Initialize match_count
    match_count = 0

    // Loop until a match is found
    do{
        if state.fetch_next_right is true{

            / If right chunk is exhausted,
            // move to the next chunk on the right
            state.left_tuple = 0
            state.right_tuple = 0
            state.fetch_next_right = false

            // Check if there are more right conditions
            if some right conditions exist:
                // If conditions exist,
                // scan the right payload data
                gstate.right_payload_data.Scan(state.right_payload)

                else:
                    // If all right conditions are exhausted,
                    // move to the next left chunk
                    state.fetch_next_left = true

                    // Handle left join: output unmatched rows
                    state.left_outer.ConstructLeftJoinResult(input, chunk)

                    return NEED_MORE_INPUT
            }
            if state.fetch_next_left is true{

                // Resolve the left condition for the current chunk
                state.lhs_executor.Execute(input, state.left_condition)

                // Reset tuples and scan conditions on the right side
                state.left_tuple = 0
                state.right_tuple = 0

                gstate.right_payload_data.Scan(state.right_payload)
                state.fetch_next_left = false
            }

            // Now perform the actual join
            // between the left and right chunks
            left_chunk = input
            right_payload = state.right_payload

            // Verify chunks
            left_chunk.Verify()
            right_payload.Verify()

            // Perform the actual join
            // using the left and right chunks
            lvector = SelectionVector(STANDARD_VECTOR_SIZE)
            rvector = SelectionVector(STANDARD_VECTOR_SIZE)

            match_count = Perform(state.left_tuple,
                                  state.right_tuple,
                                  state.left_condition,
                                  right_condition,
                                  lvector, rvector)

            // If matches are found
            if match_count > 0{

                // Set matches for both left and right sides
                state.left_outer.SetMatches(lvector,
                                             match_count)
                gstate.right_outer.SetMatches(rvector,
                                              match_count)

                // Slice the input chunks and store
                chunk.Slice(input, lvector, match_count)
                chunk.Slice(right_payload, rvector, match_count)
            }

            // If no matches are found in this iteration,
            // continue to the next iteration
            if state.right_tuple >= right_condition.size():
                state.fetch_next_right = true

            // Continue until a match is found
        } while (match_count == 0)

        // Return that more output is available

        return HAVE_MORE_OUTPUT
    }
}


```

Demonstration

```
SELECT stud.sid, enrol.cid
FROM stud JOIN enrol
ON stud.sid = enrol.sid;
```

Table Stud		
sid	sname	sage
101	A	25
102	B	26
103	A	27
104	B	23
105	A	30
107	D	30
106	C	25

Table Enrol		
sid	cid	
101	1	
101	2	
102	3	
103	3	
102	2	
105	1	
108	12	



The choice of left and right tables is done by DuckDB based on approximate cardinality estimation.




New GroupJoin Operator (GROUPJOIN_GROUPBY)


- When we find a ‘Join’ operator as a child under the ‘Group By’ node in the logical plan tree of the query.
- The Join operator can be any logical operator or physical operator, not necessarily AMUS_JOIN.
- The logical plan generated is based on the logical operators, which has to be replaced with actual physical operators.
- To implement this, whenever we see a **LOGICAL_GROUP_BY** operator, we check to see if there is a **LOGICAL_JOIN** operator somewhere in the children node.
- If such a child node exists, we replace the parent with our **GROUPJOIN_GROUP_BY** physical operator and return the generated plan.
- For this, we implemented the interface for our 'GroupJoin' physical operator (i.e. GROUPJOIN_GROUP_BY).
- For aggregation, it uses hashing similar to what is being done in the currently existing 'Group By' aggregation logic.



New GroupJoin Operator (GROUPJOIN) -- WIP



After replacing the
Group – Project - Join
With
GroupJoin Operator
(in Physical Plan)



● DuckDB New GroupJoin Operator (GROUPJOIN) -- WIP

We want to combine the Join-Project-GroupBy logic into one single operator (GROUP_JOIN)

GROUPJOIN operator will *directly* read in data from tables, do the necessary joining, and subsequent projection, so that group-by can be done

Our **GROUPJOIN_GROUPBY** can be used for [aggregation](#).

We want to extend the logic of [*GetData\(\)*](#) method in **PhysicalGroupJoin** operator

```
SourceResultType PhysicalGroupJoin::GetData(ExecutionContext &context, DataChunk &chunk,
                                         OperatorSourceInput &input) const {

    // Set the source and sink interfaces
    // Read in Data from Tables
    // Do the in-memory join
    // Set the data to sink
    return SourceResultType::FINISHED;
}
```



New GroupJoin Operator (GROUPJOIN) -- WIP

```
bool canReplaceByGroupJoin(LogicalOperator &op){  
  
    // If not a groupby, nothing to replace  
    if(op.type != LogicalOperatorType::LOGICAL_AGGREGATE_AND_GROUP_BY) return false;  
    auto &groupby = op.Cast<LogicalAggregate>();  
  
    // Check if there are groups and has a join as a child  
    if(groupby.groups.size() > 0 && groupby.children[0]->children[0]->type == LogicalOperatorType::LOGICAL_COMPARISON_JOIN){  
        return true;  
    }  
  
    return false;  
}
```

- Basic conditions to replace the Group-Project-Join with **PhysicalGroupJoin** Operator
- Additional checks based on the join conditions and group by attributes added in **PhysicalGroupJoin::IsSupported()**

Pseudocode for Plan Generator Function

```

1 unique_ptr<PhysicalOperator> PhysicalPlanGenerator::  

2     PlanGroupJoin(LogicalAggregate &op) {  

3         // Visit the children  

4         auto &join = op.children[0]->children[0]->Cast<  

5             LogicalComparisonJoin>();  

6         idx_t lhs_cardinality = join.children[0]->  

7             EstimateCardinality(context);  

8         idx_t rhs_cardinality = join.children[1]->  

9             EstimateCardinality(context);  

10        auto left = CreatePlan(*join.children[0]);  

11        auto right = CreatePlan(*join.children[1]);  

12        left->estimated_cardinality = lhs_cardinality;  

13        right->estimated_cardinality = rhs_cardinality;  

14        D_ASSERT(left && right);  

15  

16        if (join.conditions.empty()) {  

17            // No conditions: insert a cross product  

18            return make_uniq<PhysicalCrossProduct>(op.types, std  

19                ::move(left), std::move(right), op.  

20                estimated_cardinality);  

21        }  

22        unique_ptr<PhysicalOperator> plan;  

23  

24        for (auto &cond : join.conditions) {  

25            RewriteJoinCondition(*cond.right, left->types.size()  

26                );  

27        }  

28        auto condition = JoinCondition::CreateExpression(std::  

29            move(join.conditions));  

30        std::cout << "Group Join Everytime" << std::endl;  

31        // Pass the grouping and aggregate expressions  

32        plan = make_uniq<PhysicalGroupJoin>(op, std::move(left),  

33            std::move(right), std::move(condition),  

34            join.join_type, op.estimated_cardinality,  

35            op.groups, op.expressions);  

36        return plan;  

37    }  

38 } // namespace duckdb

```

DuckDB GroupJoin--WIP

Pseudocode for GetData() Function

```

1 SourceResultType PhysicalGroupJoin::  

2     GetData(ExecutionContext &context,  

3             DataChunk &chunk,  

4             OperatorSourceInput &input) {  

5  

6     // Set the source and sink interfaces  

7     // Read in Data from Tables  

8     // Do the in-memory join  

9     // Set the data to sink  

10    return SourceResultType::FINISHED;  

11 }

```

Pseudocode for checker Function

```

1 bool canReplaceByGroupJoin(LogicalOperator &op){  

2     // If not a groupby, nothing to replace  

3     if(op.type != LogicalOperatorType::  

4         LOGICAL_AGGREGATE_AND_GROUP_BY) return false;  

5     auto &groupby = op.Cast<LogicalAggregate>();  

6     // Check if there are groups and has a join as a  

7     // child  

8  

9     if(groupby.groups.size() > 0 && groupby.children  

10        [0]->children[0]->type == LogicalOperatorType::  

11        LOGICAL_COMPARISON_JOIN){  

12        return true;  

13    }  

14    return false;  

15 }

```



Tests and Experiments

To verify the implementations of **AM_US_JOIN** and **GROUPJOIN_GROUPBY aggregation only** for their correctness,

1

We created and used a small database consisting of two relations each of cardinality 8, to keep demonstration simple, as shown:

1. **Stud (Sid: INT, Sname: VARCHAR, Sage: INT)**
2. **Enrol (Sid:INT, Cid: INT)**

Table Stud			Table Enrol	
sid	sname	sage	sid	cid
101	A	25	101	1
102	B	26	101	2
103	A	27	102	3
104	B	23	103	3
105	A	30	102	2
107	D	30	105	1
106	C	25	108	12

2

We also created a database with 3 big tables each with cardinality 75, to test out on larger dataset:

- **users (user_id, first_name, last_name, address, email);**
- **products (product_id, product_name, description, price);**
- **orders (order_id, user_id, product_ordered, total_paid);**

+

We evaluated the results of these 4 queries producing ~ 5180 , ~ 383320 , ~ 70 , ~ 5625 rows respectively.
We found results consistent with the expected results

The two databases reside in *small.db* and *big.db* respectively in *\$ROOT/myduckdb/sql_files*

Summary



- reviewed the ecosystem and interfaces provided by open-source DuckDB ✓
- explored the implementation of a simple nested loop join, calling it AM_US_JOIN ✓
- tweaked the query planner and optimizer to pick our version of the Join ✓
- explored the implementation of aggregation using GROUPJOIN_GROUPBY (simple version) ✓

Future Work



- explored how we can efficiently combine both 'Join' and 'Group By' into a single 'GroupJoin' operator ✓



Thank You

Akash Maji

akashmaji@iisc.ac.in

Utkarsh Sharma

utkarsh2024@iisc.ac.in

[Link to stable version:](#)

<https://github.com/akashmaji946/myduckdb/tree/main>

