



GPU Accelerated Volume Rendering, Scalar Field Explorer

Data format and challenges

Stochastic Sampling

View Dependent Ratio

Data Importance

Phase Space Tessellation Technique

K-Buffer Concept

View Adaptive Voxelization

Optimizations

Progress So Far

Conclusion


- [1] Stochastic Volume Rendering of Multi-Phase SPH Data
- [2] High-Quality Volume Rendering of Dark Matter Simulations

Presented by:
Utkarsh Sharma
Akash Maji


Motivation: Data Format and Challenges

Given an unstructured point cloud data of scalar field, how can we do Volume Rendering ?

Where does the unstructured data come from ?



Cosmological Simulations
IllustrisTNG (Voronoi cells)




Fluid Simulations
Turbine Dataset (SPH particles)

Motivation: Data Format and Challenges

Given an unstructured data from simulations, how can we do Volume Rendering ?


Can cause artifacts due to sampling noise!



Unstructured Data

$N \times D$ co-ordinates

$N \times K$ corresponding scalars
(for data with K attributes)



Resampled Grid



$H \times W \times D \times K$ scalar values and spacing is specified

Motivation: Data Format and Challenges

The Ray Marching using the Unstructured data can use empty space skipping


Resample to a grid:

- Well defined algorithm
- Trilinear interpolation is extremely fast
- Slow Pre-processing and high memory usage
- Faster methods have less accuracy



Directly use unstructured data

- Can skip Empty/Less Important Space
- More accurate interpolation
- Slow interpolation on the fly
- Might access many points for interpolation





Assumption: SPH Data


Smoothed-particle hydrodynamics (SPH) is a computational method used for simulating the mechanics of continuum media (solid mechanics/ fluid flows). SPH treats a fluid as particles with physical attributes that contributes to a continuous field via a **kernel function** (a smooth weighting function).

$$\rho(x) = \sum_i m_i W(\|x_i - x\|, h)$$


where m_i is particle mass, W is smoothing kernel and h is it's radius.




Key Method: Stochastic Sampling



Without subsampling



Stochastic Sampling



Stratified Sampling

Use defined **sampling ratio interval**: $[p_{min}, p_{max}]$ (Different for different use cases and datasets)

The following criteria affect the probability:

- **View dependent** ratio $p_t(x)$ in real time during rendering
- **A data importance** $p_i(x)$ that is precomputed for SPH dataset

If an attribute approximation is performed during ray marching, the two criteria $p_t, p_i \in [0, 1]$ are combined multiplicatively to obtain the particle sampling probability as:

$$p(x) = p_{min} + (p_{max} - p_{min})p_t(x)p_i(x)$$


$n = p(x)|\mathcal{P}(x)|$ particles will be sampled from each grid cell for stratified

Key Method: Stochastic Sampling

View Dependent Ratio ($p_t(x)$):

Idea: Sampling points that are less visible have less contribution to the final pixel color therefore set $p_t(x)$ according to $T(x_0, x)$ for all SPH approximations at a sampling point.

To maintain a near-correct sampling at points that are almost completely visible and increase the subsampling at mostly occluded points we use a quadratic falloff as:



Transmittance between points x_0 and x_1 is given by $T(x_0, x_i)$

Occlusion can be obtained by $(1 - T(x_0, x_i))$

Sample less from the occluded region and more from visible region


Key Method: Stochastic Sampling

Precomputed data importance ($p_i(x)$):

Idea: If neighboring particles have a higher attribute variation, the variance in the sampling process is increased and more particles should be considered in the SPH approximation

First, all attribute values of the particles in $\mathcal{P}(x)$ are binned in a normalized histogram $h(\mathbf{b})$ with N_{bins} bins. Then, the entropy is obtained as:

$$\mathbb{H}_{\mathcal{P}}(\mathbf{A}) = - \sum_{b=0}^{N_{bins}-1} h(b) \log_2(h(b))$$



Key Method: Stochastic Sampling

Precomputed data importance ($p_i(x)$):


Idea: If neighboring particles have a higher attribute variation, the variance in the sampling process is increased and more particles should be considered in the SPH approximation

First, all attribute values of the particles in $\mathcal{P}(x)$ are binned in a normalized histogram $h(b)$ with N_{bins} bins. Then, the entropy is obtained as:

$$\mathbb{H}_{\mathcal{P}(x)}(A) = - \sum_{b=0}^{N_{bins}-1} h(b) \log_2(h(b))$$

Normalize the values to be independent from the histogram size:

$$\bar{\mathbb{H}}_{\mathcal{P}(x)}(A) = \frac{2^{\mathbb{H}_{\mathcal{P}(x)}}}{N_{bins}}$$



Define the importance value for position x , with set of particles in the local neighborhood $\mathcal{P}(x)$, as:

$$p_i(\mathbf{x}) = \begin{cases} 0, & \text{if } |\mathcal{P}(\mathbf{x})| = 0 \\ 1, & \text{if } \exists i, j \in \mathcal{P}(\mathbf{x}) : \pi_i \neq \pi_j, \\ \bar{\mathbb{H}}_{\mathcal{P}(\mathbf{x})}(A), & \text{otherwise} \end{cases}$$

Results




Figure 8: Equal-time rendering using uniform particle sampling compared to our non-uniform sampling strategy in Figure 3(b). While both images took about 5.3 s to render, our strategy has significantly lower error.




Figure 9: Rendering using volume resampling (269.896 s pre-processing, 0.293 s rendering) with inlet comparisons to our high quality stochastic sampling configuration (0.029 s, 6.286 s) and a rendering without any sampling (0 s, 46.612 s).

Dataset	Particles	Preproc.	Fast Sampling		Quality Sampling		No Sampling	Volume Resampling	
			Frame Time	RMSE	Frame Time	RMSE		Preproc.	Frame Time
Bubble	4,347,225	10 ms	0.993 s	8.277	5.050 s	4.360	10.274 s	143.234 s	124 ms
Spray Nozzle	43,188,662	28 ms	0.982 s	23.108	8.350 s	5.391	44.895 s	447.261 s	82 ms
Turbine	86,473,832	29 ms	0.962 s	23.828	5.347 s	7.187	39.624 s	269.896 s	293 ms

Root-mean-square errors are measured compared to the rendering without stochastic sampling.

High Quality Volume Rendering of Dark Matter Simulations

N-body dark matter simulations are essential for studying the large-scale structure of the Universe

We rely on accurate **mass densities** for identifying features of the **cosmic web**

Traditional Approaches

Problem: suffer from **artifacts due to sampling noise**

Alternate Approach

The **Phase Space Element Approach (PSEA)** constructs a tessellation of the dark matter sheet (3D) embedded in **6-dimensional phase space**.

This approach yields **density fields of very high quality** by summing the density contributions of all overlapping elements in **configuration space**

Paper Goal: To introduce a robust volume rendering approach that achieves **high image quality** at **interactive** frame rate for PSEA

So, we want PESA, but...

The PSEA generates depth-complex **tetrahedral meshes**

They suffer from an **excessive amount of self-intersections** due to dark matter sheet folding over itself

In **overdense regions**, it is common for up to **10^4 cells to overlap** a single spatial location

Previous Attempts:

They suffered from **image-noise**; used a **binning approach** for high depth complexity areas

The Novel Hybrid Volume Rendering Algorithm [4 Steps] for PESA


It combines a **single-pass k-buffer approach** with a **view-adaptive voxelization grid**

It achieves **Order-Independent Transparency (OIT)** and guarantee correct depth sorting for the **k closest**, highly overlapping fragments

Needs:


Leverages atomic operations on 64-bit integers on the GPU

Parameter choices (k , voxel resolution) affect quality and memory




Part-1: Geometry Construction

- Particle positions (vertices) are uploaded to GPU(3D-texture) and accessed via Lagrangian grid indices.
- GS constructs the **six tetrahedra** per cell and calculates their mass densities.
- FS computes **entry point depth (d{in})** and the **intersection length(dt)** inside tetrahedron for the viewing ray.




Part-2: K-Buffer Construction



- Handle extreme depth complexity **near camera** by creating a perfectly **depth-sorted** list of the first k fragments.
- The **data item** is inserted into the pixel's fragment list by looping over the k -entries, atomically exchanging it
- **Result:** depth-sorted list of the viewing rays' entry and exit points for the **k -closest** tetrahedra faces
(along with the tetrahedra's density ρ at entry points and $-\rho$ at the exit points)

Part-3: View Adaptive Voxelization

- Capture the density contributions of tetrahedra beyond k -buffer coverage
- [Viewport-aligned 3D-texture](#); **logarithmic depth spacing**.
- Density contributions (ρ_T) are [atomically added](#) and weighted by (d_T)
- Prevents [high-frequency image noise](#) because it correctly accumulates the contributions of **thin** tetrahedra (not possible using binning)



Part-4: Colour-mapping and Blending

- Combine foreground (k -buffer) and background (voxel grid) density into the final image color.
- Densities are combined in a **front-to-back, view-consistent order**.
- Render by looping over the entries in the k -buffer, accumulate the density contributions, map the current density at each interval to colors and opacities.
- Analogously, the densities stored in the view-adaptive voxelization grid are sampled, mapped, blended.


Optimization1: Tiled Rendering & View-Frustum Culling

Generating geometry in the **GS** is a major overhead, especially for outside the view

Solution:

Subdivide the data into cubical blocks (e.g., 8^3 cells per block)

Only blocks whose **AABB** intersect the current tile's view frustum are rendered



Optimization2: View-Dependent Multiresolution (LOD)

Rendering full resolution data is slow

Solution:


Resolution is determined by the **screen-space extension** of the intersecting blocks' **AABB**.

Data for lower resolutions are generated by **downsampling**.


Results

- Achieves **high quality**; optimizations give **interactive** frame rates.
- Captures subtle features: **caustic structures**
- Single-pass **k-buffer** + View-adaptive **voxelization**: resolves depth complexity and image noise.

Results



Cell Projection Approach



Our Hybrid Approach



Progress

What we Have


- CUDA + Python based Volume renderer and Transfer Function Editor
- Support for IllustrisTNG (cosmological data) and regular grid
- Simple GUI to load and visualize the data
- A simple Neural Representation and renderer in python (incomplete)

What we are working on

- Support for data importance sampling and empty space skipping
- Support for different/compressed representation of volume data
- Quality comparison and metrics for comparison




Comparison of Different
Resampling Techniques
Illustris Subhalo Cutout 568590



Neural
Representation
(Incomplete)

A visualization of the Head dataset
(open sci vis) to verify the
implementation





Thank You

References

- M. Piochowiak, T. Rapp, and C. Dachsbacher, “**Stochastic Volume Rendering of Multi-Phase SPH Data**,” *Computer Graphics Forum*, vol. 00, no. 00, pp. 1–13, 2020. doi: 10.1111/cgf.14121.
- Ralf Kaehler, "High-Quality Volume Rendering of Dark Matter Simulations" in Proc. IS&T Int'l. Symp. on Electronic Imaging: Visualization and Data Analysis, 2018, pp 333-1 - 333-8, <https://doi.org/10.2352/ISSN.2470-1173.2018.01.VDA-333>

Presented by:
Utkarsh Sharma
Akash Maji